



Database design

- Entity-Relationship Model
- Conceptual design
- Logical design
- Normalization

Database design

0



1

1

Entity-Relationship Model

- Life cycle of an information system
- Database design
- Entities and Relationships
- Attributes
- Identifiers
- Generalization
- Documenting E-R Schematics
- UML and E-R

Entity-Relationship Model

Database design



2

2



3

3

Database design

- The design of a database is one of the activities of the process of developing an information system
 - must be seen in the broader context of the life cycle of an information system

Life cycle of an information system

Database design



4

4



5

5

Life cycle of an information system

Determination of the costs of the different alternatives and the priorities for the implementation of each system component

Feasibility study

DBGI

6

Life cycle of an information system

- Definition of the properties and functionalities of the information system
- Requires user interaction
- Produces a comprehensive, but informal, description of the system to be implemented

Feasibility study

Collection and analysis of the requirements

DBGI

7

Life cycle of an information system

- Divided into data and application design
- Produces formal descriptions

Feasibility study

Collection and analysis of the requirements

Design

DBGI

8

Life cycle of an information system

- Implementation of the information system according to the characteristics defined in the design phase

Feasibility study

Collection and analysis of the requirements

Design

Implementation

DBGI

9

Life cycle of an information system

- Verification of the correct functioning and quality of the information system
- It can lead to changes in requirements or design revision

Feasibility study

Collection and analysis of the requirements

Design

Implementation

Validation and testing

DBGI

10

Life cycle of an information system

- System operation
- Requires maintenance and managing operations

Feasibility study

Collection and analysis of the requirements

Design

Implementation

Validation and testing

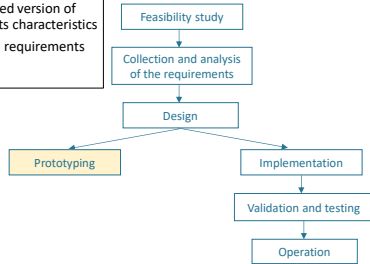
Operation

DBGI

11

Life cycle of an information system


- Quickly create a simplified version of the system to evaluate its characteristics
- It can lead to changes in requirements or design revision



D&G

12

Life cycle of an information system



D&G

13

Database design

Database design

D&G

14

Database design

- The database is an important component of the entire system
- Data-driven design methodology
 - the design of the database precedes that of the applications that use it
 - greater attention to the design phase than to the other phases

D&G

15

Design Methodology

- A design methodology consists of
 - decomposition of the project activity into successive and independent phases
 - strategies to be followed in the various phases and criteria for choosing the best strategy
 - reference models to describe the input and output data of the various phases

D&G

16

Properties of the methodology

Generality

- can be used regardless of the problem and the tools available

Quality of the result

- in terms of correctness, completeness and efficiency with respect to the resources used

Ease of use

- of both strategies and reference models

D&G

17

Data-driven design

- For databases, methodology based on separating two key decisions
 - what** to represent in the database
 - conceptual design
 - how** to represent it
 - logical and physical design

DBGI

18

18

Stages of database design



Informal specification of the reality of interest


- Application properties
- Application functionalities

DBGI

19

19

Stages of database design



Representation of informal specifications in the form of a **conceptual diagram**


- formal and complete description, referring to a conceptual model
- Independent from implementation aspects (data model)
- the aim is to represent the **information content** of the database

DBGI

20

20

Stages of database design



Translating the conceptual schema into the **logical schema**


- depends on the chosen data **logic model**
- takes into account the **optimization** of data processing operation
- schema quality** verified by formal techniques (normalization)

DBGI

21

21

Stages of database design



Specification of the **physical data storage parameters** (File and index organization)

- produces a **physical model**, which depends on the chosen DBMS

DBGI

22

22

Entity-Relationship Model

Database design

DBGI

23

23

The E-R (Entity-Relationship) model

- It is the most widely used conceptual model
- Provides constructs to describe specifications about data structure
 - in a simple and understandable way
 - with a graphic formalism
 - independent of the data model, which can be chosen later
- Several variants are available

Main elements of the E-R model

- Entity
- Relations
- Attributes
- Identifiers
- Generalizations and subsets

DBGI

24

24

DBGI

25

25

Entity


Entity name

- It represents classes of real-world objects (people, things, events, ...), which have
 - common Properties
 - autonomous existence
- Examples: Employee, Student, Article
- An occurrence of an entity is an object of the class that the entity represents

26

26

Example of entities



DBGI

27

27

Relationship




- Represents a logical link between two or more entities
- Examples: exam between student and course, residence between person and municipality
- Not to be confused with the relation of the relational model
 - sometimes referred to as association

28

28

Relationship examples




DBGI

29

29

Occurrences of a relationship

- An occurrence of a relationship is an n-tuple (pair in the case of a binary relationship) consisting of occurrences of entities, one for each of the entities involved
- There can be no identical n-tuples



D.B.G.

30

30

Cardinality of binary relationships

- They are specified for each entity that participates in a relationship
- Describe the minimum and maximum number of occurrences of a relationship in which an occurrence of an entity can participate
 - minimum** can be either
 - 0 (optional participation)
 - 1 (participation required)
 - maximum** varies between
 - 1 (at most one occurrence)
 - N (arbitrary number of occurrences)


D.B.G.

31

31

Cardinality of binary relationships

- 1-to-1 relationship




D.B.G.

32

32

Cardinality of binary relationships

- 1-to-N (many) relationship




D.B.G.

33

33

Cardinality of binary relationships

- N-to-N (Many to Many) relationship




D.B.G.

34

34

Limitations of binary relationships



- It is not possible for a student to take the same exam more than once


D.B.G.

35

35

Ternary relationship

- A student may take the same exam more than once at different times.




- Example of an instance of the EXAM report
- | | | |
|-------|-------|-------|
| s_1 | c_1 | t_1 |
| s_1 | c_1 | t_2 |
| ... | | |

D8G

36

36

Occurrences of a ternary relationship




D8G

37

37

Occurrences of a ternary relationship




D8G

38

38

Occurrences of a ternary relationship




D8G

39

39

Occurrences of a ternary relationship




D8G

40

40

Cardinality of ternary relationships




- Minimum cardinalities are rarely 1 for all entities involved in a relationship
- The maximum cardinalities of an n-ary relationship are (practically) always N
 - if the participation of an entity E has a maximum cardinality of 1, it is possible to eliminate the n-ary relationship and associate the entity E with the others by binary relations

D8G

41

41

Recursive relationship




- Relationship between an entity and itself
- If the relationship is not symmetrical, the two roles of the entity must be defined

DBGi

42

42

Recursive relationship




DBGi

43

43

Recursive relationship




- An employee might have several managers

DBGi

44

44

Attribute




- Describes an elementary property of an entity or relationship
- Examples
 - surname, first name, student ID are attributes that describe the student entity
 - grade is an attribute that describes the exam relationship
- Each attribute is characterized by the **domain**, the set of admissible values for the attribute

45

45

Examples of attributes




DBGi

46

46

Composite attribute



- Group of attributes that have closely connected meanings or uses.
- Example:




DBGi

47

47

Cardinality of an attribute

- Can be specified for entity or relationship attributes
- Describes the minimum and maximum number of attribute values associated with an occurrence of an entity or relationship
 - if it is omitted it corresponds to (1,1)
 - minimum 0** corresponds to an attribute that admits a null value
 - maximum N** corresponds to an attribute that can have more than one value for the same occurrence (multivalued attribute)



48

D.B.G.

48

Identifier

- It is specified for each entity
- Describes the concepts (attributes and/or entities) of the schema that allow you to uniquely identify the occurrences of the entities
 - each entity must have at least one identifier
 - there can be more than one appropriate identifier for an entity
- The identifier can be
 - internal or external
 - simple or composite


49

49

Internal identifier


Simple

- consisting of a single attribute



Composite

- consisting of multiple attributes




50

D.B.G.

50

External identifier

- An entity that does not have internal attributes sufficient to define an identifier is called a **weak entity**
- The weak entity must participate with cardinality (1,1) in each of the relationships that provide part of the identifier




51

51

Remarks

- An external identifier can involve an entity that is itself externally identified
- No identification cycles should be generated




52

D.B.G.

52

Remarks

- Relationships do **not** have identifiers




53

53

Generalization


It describes a logical link between an entity E and one or more entities E_1, E_2, \dots, E_n that are particular cases of E .

- E is called **parent entity**, and is a **generalization** of E_1, E_2, \dots, E_n
- E_1, E_2, \dots, E_n are called **child entities**, and are **specializations** of E



54

Generalization: example




55

54

55


Generalization: example



56

56

Generalization: example



57

57


Generalization: example



58

58


Generalization: example



59

59

Generalization: example



D.B.G.

60

60

Generalization: properties

- Each occurrence of a child entity is also an occurrence of the parent entity
- Each property of the parent entity (attributes, identifiers, relationships, other generalizations) is also a property of each child entity
 - property known as **inheritance**
- An entity can be involved in multiple different generalizations

D.B.G.

61

61

Generalization: properties


- Orthogonal characteristics
 - total** generalization if each instance of the parent entity is an instance of at least one of the child entities, **partial** otherwise.
 - exclusive** if each instance of the parent entity is at most one instance of one of the child entities, **overlapping** otherwise.

D.B.G.

62

62

Generalization: incorrect example




D.B.G.

63

63

Generalization: incorrect example




D.B.G.

64

64

Generalization: incorrect example




D.B.G.

65

65

Generalization: incorrect example




D.B.G.

66

66

Generalization: correct example




D.B.G.

67

67

Subset

- Particular case of generalization with only one child entity
- the generalization is always partial and exclusive



D.B.G.

68

68

ER Model Documentation


Database design

D.B.G.

69

69

Documenting E-R models



D.B.G.

70

70

Documenting E-R models



Data Dictionary

Enrich the E-R schema with natural language descriptions of entities, relationships, and attributes

D.B.G.

71

71

Data dictionary: example

Entity	Description	Attributes	Identifier
Student	University student	Student ID, Surname, Name, CFU acquired, Grades average	Student ID
Professor	University professor	Professor ID, Department, Surname, Name	Professor ID
Course	Courses offered by the university	Course code, Name, CFU	Course code
Time	Dates on which exams were taken	Date	Date

DBGi

72

72

Data dictionary: example

Relationship	Description	Entities involved	Attributes
Exam	It associates a student to the exams taken and memorizes the mark obtained	Student (0,N), Course (0,N), Time (1,N)	Grade
Holder	It associates each course to the professor who teaches the course.	Course (1,1), Professor (0,N)	

DBGi

73

73

Documenting E-R models



Data Dictionary

Enrich the E-R schema with natural language descriptions of entities, relationships, and attributes



Data Integrity Constraints

They cannot always be explicitly stated in an E-R scheme
They can be described in natural language

DBGi

74

74

Data integrity constraints: examples

Integrity constraints	
RV1	The grade of an exam can only take values between 0 and 30
RV2	Each student cannot pass the same exam twice
RV3	A student may not take more than three exams for the same course during the same academic year

DBGi

75

75

Documenting E-R models



Data Dictionary

Enrich the E-R schema with natural language descriptions of entities, relationships, and attributes



Data Integrity Constraints

They cannot always be explicitly stated in an E-R scheme
They can be described in natural language



Data Derivation Rules

Explicitly define that a concept of the schema can be obtained (by inference or arithmetic calculation) from other concepts of the schema

DBGi

76

76

Derivation rules: examples

Derivation rules	
RD1	The number of credits acquired by a student is obtained by adding the number of credits of the courses for which the student has passed the exam
RD2	The average mark is obtained by calculating the average of the marks of the exams passed by a student

DBGi

77

77

UML vs ER

Database design

DBG

78

78

UML and ER

UML (Unified Modeling Language)

- Modeling a software application
 - structural and behavioural aspects (data, operations, processes and architectures)
- Rich formalism
 - class diagram, actor diagram, sequence diagram, communication diagram, state diagram,...

ER

- Modeling a database
 - structural aspects of an application
 - elements tailored to the modelling of a database

DBG

79

79

UML vs ER

- Different formalisms
 - The class diagram of an application is different from the E-R schema of the database
 - The class diagram, even if designed for different uses, can be adapted for the description of the conceptual design of a database
- Main Differences of UML vs ER
 - no standard notation to define identifiers
 - ability to add notes to comment on diagrams
 - possibility to indicate the direction of navigation of an association (not relevant in the design of a database)

80

80

Politecnico di Torino

DBG

Examples of time representation in the ER scheme

Database design

0

Time in E-R models

- Time is needed to represent:
 - Events
 - Temporal changes in values and/or relationships
- Time can be modelled using:
 - Temporal attributes
 - Binary relationships
 - Ternary relationships
 - Entity historicization

1

Temporal attributes

- Temporal information related to a single entity or relationship
- Unique events for each entity instance
 - Example: birth date, film release date

2

Binary relationship

- Example: A set of sensors are available, each identified by a unique code and present within a building.
 - It is requested store the different temperature values detected by each sensor at **different time points**.

3

Binary relationship

- We want to represent a time series of events related to an Entity E of the ER diagram
- The information of interest is represented by introducing
 - a **Time** entity
 - identified by time information about **when** the event occurs/starts (e.g. timestamp, date, date and time)
 - a **binary relationship R** that connects the **Time** entity to the **E** entity
 - The (if any) information on the duration and/or time of the end of the event and/or on other aspects that characterize the occurrence of the event at different time points are attributes of the relationship **R**

4

Workaround: Weak Entity


Example: A set of sensors are available, each identified by a unique code and present within a building.

- It is requested store the different temperature values detected by each sensor at **different time points**.

5

Workaround: Weak Entity

- The Event Entity **E**
 - is a weak entity identified internally by the set of attributes that represent the instant in which the event itself begins/occurs
 - the characteristics of the event are attributed directly to the weak entity **E**



D&G


6

6

Ternary relationship

Example: It is requested to store the product purchases made by each customer.

- Each customer is uniquely identified by their tax code.
- Each product is uniquely identified by the barcode and characterized by the name.
- Suppose each customer can buy the same product at different times of the same day.



D&G

7

7

Ternary relationship

- We want to represent a time series of events expressed through an association/relationship between two Entities **E1** and **E2** of the ER diagram
- The information of interest is represented by introducing
 - a **Time** entity
 - identified by time information about **when** the event occurs/starts (e.g. timestamp, date, date and time)
 - a **ternary relationship R** that connects the **Time** entity, the **E1** entity, and the **E2** entity
 - The (if any) information on the **duration** and/or time of the **end** of the event and/or on other aspects that characterize the occurrence of the event at different time points are attributes of the relationship **R**

D&G


8

8

Ternary Relationship

Example: A company that provides IT consulting wants to store the work done by its employees for each project.

- Each project is identified by an alphanumeric code and characterized by a title.
- Employees who work at the company are identified by their tax code and characterized by their first and last name and email.
- It is requested to store the time periods (start date, end date) in which an employee works on a project. Multiple employees can work at the same time on the same project.



D&G

9

9

Historicized Entity

Example: It is requested to store the lectures given by each teacher for each course.

- Each teacher is uniquely identified by a numeric ID and is characterized by surname and first name.
- Each course is identified by an alphanumeric code and characterized by its name.
- Each lesson is characterized by the date and time slot (start time and end time) in which it is held and by the course for which the lesson is delivered. Suppose that each teacher can deliver a maximum of one lecture in the same time slot.

D&G

10

10

Historicized entity


- You want to represent
 - An event that involves two entities
 - There are **constraints** on an entity's participation in multiple events
- The information of interest is represented by introducing
 - a weak entity **E** identified internally by the time information about **when** the event occurs/starts (e.g. timestamp, date, date and time) and externally by the relationship linked to the entity that **cannot** participate in two events at the same time
 - the characteristics of the event are attributed directly to the weak entity **E**
 - entity **E** participates with cardinality (1,1) in relationships with other entities

D&G

11

11

Historicized entity



DBGI

12

12

Recap

Binary



A student can take the exam for each course only once.

Translation:

STUDENT(StudentID)
COURSE(CourseID)
Exam(StudentID, CourseID, Date)

Ternary




A student can take the exam for each course multiple times on different dates.

Translation:

STUDENT(StudentID)
COURSE(CourseID)
~~TIME(Date)~~ Can be omitted
Exam(StudentID, CourseID, Date)

Historicized



A student can take the exam for each course multiple times. A student can only take one exam on a given date.


Translation:

STUDENT(StudentID)
COURSE(CourseID)
Exam(StudentID, Date, CourseID)


DBGI

13

13



Politecnico di Torino




Logical Design

Database Design

0

Logical Design (1/2)

- Introduction
- Restructuring of the Entity-Relationship schema
- Removing generalizations
- Partitioning of concepts
- Removing multivalued attributes
- Removing composed attributes
- Selection of primary identifiers




1

1

Logical Design (2/2)

- Translation into the relational model
 - entity and many-to-many relationships
 - one-to-many relationships
 - one-to-one relationships
 - entities with external identifiers
 - ternary relationships




2

2

Logical Design


Introduction



3


3

Logical design




Selection of the logical model

Relational model




Objective

Definition of a relational logical schema corresponding to the starting ER schema



Important

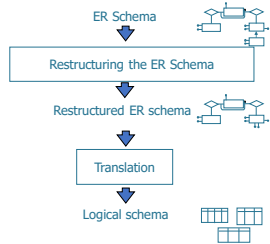

Simplification of the ER schema to make it representable by the relational model
Optimization to increase the efficiency of queries



4

4

Logical design steps

5

5

Translation to the relational model

entities and many-to-many relationships



6

Translation to the relational model

- It is executed on the restructured ER schema
 - i.e., the schema without hierarchies, multivalued attributes and composite attributes
- Transformations
 - Each entity is translated into a table with the same attributes
 - For relationships we need to consider the maximum cardinality



7

Entity Translation

Translating the ER Schema into the Relational Model



8

Entity Translation

- Each entity corresponds to a table with the same attributes
 - the **attributes** of the entity constitute the **schema** of the table
- The identifier (simple or composite) of the Entity becomes the primary key of the table
- Optional Entity attributes are attributes that can be NULL
 - They are highlighted with "*" in the table schema




9

Entity

Conceptual model

Logical model



Person(SSN, Name, Surname, Profession*)

- Underlined primary key
- Optional attributes indicated with an asterisk



10

Relationship translation

Translating the ER Schema into the Relational Logic Model



11

Relationship translation

- To translate a relationship
 - Step 1: The Entities participating in the Relationship are first translated
 - Step 2: The Relationship is then translated
 - Different translation rules for binary and ternary Relationships
- For a Binary Relationship, it is necessary to consider the maximum and minimum **cardinality** with which the Entities participate in the Relationship

DBG

12

12

Translation of Binary Relationships

Translating the ER Schema into the Relational Model

DBG

13

13

Many-to-many binary relationship

Conceptual model



Logical model

Student(StudentID, Name, Surname)
 Course(CodC, Name)
 Exam(StudentID, CodC, Grade)

- Each many-to-many relationship corresponds to a table
 - The primary key is the combination of the identifiers of the two linked entities
 - The attributes of the table that corresponds to the relationship can be renamed (required in case of recursive relationships)


DBG

14

14

Recursive many-to-many binary relationship

Conceptual model



Logical model

Product (P_ID, Name, Cost)
 Composed(CodCompound, CodComponent, Quantity)

- Each many-to-many relationship corresponds to a table
 - The primary key is the combination of the identifiers of the two linked entities
 - The attributes of the table that corresponds to the relationship can be renamed (required in case of recursive relationships)

DBG

15

15

One-to-many binary relationship

- Two translation modes are possible
 - by means of attributes
 - by means of a new table


DBG

16

16

One-to-Many Binary Relationship: using attributes

Conceptual model



Logical model

Person (SSN, Name, Surname, CityID, Date)
 City (CityID, Name)

- It is used when participation of the entity that participates with a maximum cardinality of 1 is mandatory (minimum cardinality of 1)


DBG

17

17

One-to-many binary relationship: using attributes or a new table

Conceptual model



Logical model

Alternative 1: Translation using attributes

Student (StudentID, Name, Surname, **NameUniv***, **Date***)
University (Name, City)

Alternative 2: Translation using a new table

Student (StudentID, Name, Surname)
University (Name, City)
Graduation (**StudentID**, **NameUniv**, **Date**)

- It is used when participation of the entity that participates with a maximum cardinality of 1 is optional (minimum cardinality of 0)

DBGi

18

18

One-to-one binary relationship

- Multiple translations are possible

- depends on the value of the minimum cardinality


DBGi

19

19

One-to-one binary relationship

Conceptual model



Logical model

Alternative 1

Rector (ID, Name, Surname, **UnivName**, **Date**)
University (Name, City)

Alternative 2

Rector (ID, Name, Surname)
University (Name, City, **ID**, **Date**)

- It is used when both entities participate with a maximum cardinality of 1 in the relationship, and participation is mandatory for both entities (minimum cardinality of 1)


DBGi

20

20

One-to-one binary relationship

Conceptual model



Logical model

Professor (ID, Name, Surname)
University (Name, City, **RectorID**, **ElectionDate**)

- It is used when both entities participate with a maximum cardinality of 1 in the relationship, but participation is mandatory only for one entities (minimum cardinality of 1)

DBGi

21

21

One-to-one binary relationship

Conceptual model



Logical model

Alternative 1

Professor (ID, Name, Surname)
University (Name, City)
Rector (**RectorID**, **UniversityName**, **ElectionDate**)

Alternative 2

Professor (ID, Name, Surname)
University (Name, City)
Rector (**RectorID**, **UniversityName**, **ElectionDate**)

Alternative 3

Professor (ID, Name, Surname)
University (Name, City, **RectorID***, **ElectionDate***)

- It is used when both entities participate with a maximum cardinality of 1 in the relationship, and participation is optional for both entities (minimum cardinality of 0)

DBGi

22

22

Translation of Ternary Relationships

Translating the ER Schema into the Relational Model


DBGi

23

23

Ternary Relationship

Conceptual model



Logical model

Student(StudentID, Name, Surname)
 Course(CodC, Name)
 Time(Date)
 Exam(StudentID, CodC, Date, Grade)

24


Translating Entities with External Identifier

Translating the ER Schema into the Relational Model

25

Entities with an external identifier

Conceptual model



Logical model

Cinema (CodC, Name, Address)
 Hall (Number, CodC, Capacity)

- The relationship is represented together with the identifier
- The relationship contributes to the definition of the weak entity identifier

26


Referential integrity constraints

Translating the ER Schema into the Relational Model

27

Referential integrity constraints

Conceptual model



Logical model

Student(StudentID, Name, Surname)
 Course(CodC, Name)
 Exam(StudentID, CodC, Grade)


- Relationships Represent Referential Integrity Constraints

Exam(StudentID) REFERENCES Student(StudentID)
 Exam(CodC) REFERENCES Course(CodC)

28

Referential integrity constraints

Conceptual model



Logical model

Room (ID)
 Bookshelf (ID, B_ID, Name)
 Shelf (ID, B_ID, Shelf_ID)


- If the referenced key consists of multiple attributes, the referential integrity constraint is imposed on the attribute set

Bookshelf(ID) REFERENCES Room(ID)
 Shelf (ID, B_ID) REFERENCES Bookshelf (ID, B_ID)

29

Referential integrity constraints

Conceptual model



Logical model

Room (ID)
Bookshelf (ID, B_ID, Name)
Shelf (ID, B_ID, Shelf_ID)

- If the referenced key consists of multiple attributes, the referential integrity constraint is imposed on the attribute set

Bookshelf(ID) REFERENCES Room(ID)

Shelf (ID) REFERENCES Bookshelf (ID)

Shelf(B_ID) REFERENCES Bookshelf (B_ID)

Wrong constraints!

DBG

30

30

Restructuring the ER model

Restructuring the ER model

DBG

31

31

Restructuring the ER model

- The restructured ER model takes into account implementation aspects
 - It is no longer a conceptual model
- Objectives
 - To eliminate constructs for which there is no direct representation in the relational model
 - To transform the data representation in order to increase the efficiency of data access operations

DBG

32

32

Restructuring tasks

- Eliminating composite attributes
- Eliminating multivalued attributes
- Eliminating generalizations
- Analysis of redundancies
- Partitioning concepts (Entities, Relationships)
- Choosing primary identifiers

DBG

33

33

Eliminating composite attributes

Restructuring the ER model

DBG

34

34

Eliminating composite attributes


- Composite attributes are not representable in the relational model.
- Attributes can be deleted by:
 - separately representing individual sub-attributes
 - if you need to access each attribute separately
 - Introducing a single attribute that represents the concatenation of the composite attributes
 - if access to the overall information is sufficient

DBG

35

35

Option 1: separate attributes




DBG

36

36

Option 2: single attribute



DBG

37

37

Eliminating multivalued attributes

Restructuring the ER model

DBG

38

38

Eliminating multivalued attributes


- They cannot be represented in the relational model
- Multivalued attributes are represented using a relationship between:
 - the initial entity
 - a new entity
- Pay attention to the cardinality of the new relationship

DBG

39

39

Shared information




A person can have more than one educational qualification and that the same educational qualification can be held by several people

DBG

40

40

Shared information: Has_qualification cardinality




DBG

41

41

Shared information: *Has_qualification* cardinality




A person can have more than one educational qualification and that the same educational qualification can be held by several people

D&G

42

42

Unique information




A person can have more than one telephone number, but a given telephone number can be held only by one person

D&G

43

43

Unique information: *Has_telephone* cardinality




D&G

44

44

Unique information: *Has_telephone* cardinality



A person can have more than one telephone number, but a given telephone number can be held only by one person

D&G

45

45

Removing generalizations

Restructuring the ER model

D&G

46

46

Removing generalizations


- The relational model does not allow direct representation of generalizations of the ER model
 - We need, therefore, to transform these into entities and relationships
- Possible methods:
 - Child entities merged into parent entity
 - Parent entity merged into child entities
 - Generalization translated into relationships

D&G

47

47

Example




D&G

48

48

Merging child entities into the parent entity




D&G

49

49

Attributes of child entities




D&G

50

50

Relationships with child entities




D&G

51

51

The «Type» attribute




- The **Type** attribute indicates the original entity (doctor or volunteer) to which each occurrence of the parent entity (staff) belongs

D&G

52

52

Merging child entities into the parent entity




- Can be used for all types of generalization
 - in case of overlapping entities, many combinations are possible as Type values, e.g., skier and sailor

D&G

53

53

Back to the example




DBG

54

54

Merging the parent into the child entities




DBG

55

55

Attributes of the parent entity




DBG

56

56

Relationships with parent entity




- Relationships with the parent entity need to be split

DBG

57

57

Cardinality of Works in relationship



DBG

58

58

Merging the parent into the child entities




- Cannot be used for **partial** generalizations
 - however, generalizations can be transformed from partial to total by adding a new entity **Others**
- Cannot be used for **overlapping** generalizations
 - due to duplicate identifiers

DBG

59

59

Back to the original example




D&G

60

60

Generalization translated into a relationship




D&G

61

61

Child entities' identifier




D&G

62

62

Cardinality of is a relationship




D&G

63

63

Generalization translated into relationships



- This solution is more general and can be used for all generalizations
- But it may be expensive to reconstruct the original data

D&G

64

64

Assessment of alternatives

- Merging child entities into parent entity is appropriate when:
 - access operations apply to instances and attributes of child and parent entities more or less in the same way (optimize data access)
 - child entities are mildly differentiated (few null values)
- Merging parent entity into child entities is appropriate when:
 - the generalization is total
 - there are operations that refer only to specific child entities and therefore it is useful to distinguish between different child entities (optimize data access)
- "Mixed" representations are also possible:
 - there are operations that refer only to instances of some child entities (optimize data access)
- In the presence of hierarchical generalization, apply the same procedure, starting from the lower levels

D&G

65

65

Redundancy analysis

Restructuring the ER model

DBG

66

66

Redundancy analysis

- They represent information that is relevant to the application, but can be derived from other concepts
 - it must be decided whether to keep them
- Effects of redundancies on the logical schema
 - simplifying and speeding up queries
 - increased complexity and slower updates
 - increased storage requirements


DBG

67

67

Redundant attribute example

- The Avg_grade attribute is redundant:
 - Useful for speeding up queries that require calculating the average of students' grades
 - if preserved, the relational schema must be supplemented with proper documentation that the attribute is redundant (and derivation rules)



DBG

68

68

Partitioning concepts

Restructuring the ER model

DBG

69

69

Partitioning of concepts


- Partitioning of entities and relationships
 - better representation of different concepts
 - separating attributes of the same concept that are accessed by different operation
 - improve the efficiency of the operations

DBG

70

70

Entity partitioning




DBG

71

71

Relationship partitioning



DBG

72

72

Choosing Primary Identifiers

Restructuring the ER model

DBG

73

73

Selection of primary identifiers

- It is necessary to define the relation *primary keys*
- The criteria for this decision are as follows
 - Attributes with **null** values **cannot** form primary identifiers.
 - Just **one** (better) or **few** attributes
 - An **internal** identifier is preferable to an external one
 - It is used by many operations to access the occurrences
- It may be useful to introduce an additional attribute to represent the entity, often called code or ID, e.g. «ProductCode»

DBG

74

74

Politecnico di Torino

DBG

Database design

Example

0

Conceptual design (1/2)

- Database design stages
- Design example: problem specifica
- Design example: main concepts
- Design example: model refinement (I)
- Design example: model refinement (II)
- Design example: model refinement (III)

DBG

1

Conceptual Design (2/2)

- Design example: representation of time (I)
- Design example: representation of time (II)
- Design example: representation of time (III)

DBG

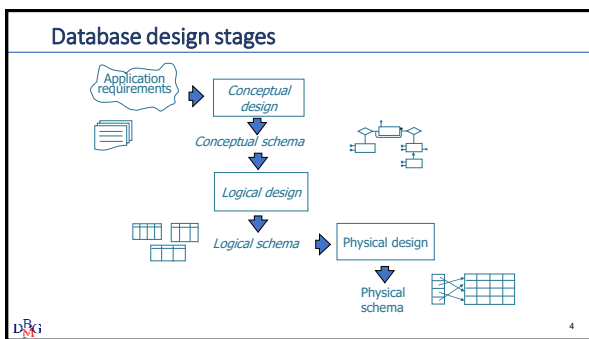
2

Database design

Database design stages

DBG

3



4

Requirements collection and analysis

- Requirements collection
 - problems that the application must solve
 - static and dynamic features of the application
- Requirements analysis
 - clarification and organization of the specifications
- Interconnected and scarcely standardizable activities

DBG

5

Requirements sources

- Application users
 - interviews
 - written documentation
- Existing documentation
 - regulations
 - internal rules
 - forms
- Existing systems
 - applications to replace or interact with



6

6

Requirements collection

- System users play an important role
 - high-level users have a more general view, but they often do not know the details
 - different users can provide different information (complementary or contradictory)



7

7

Requirements collection

- In practice:
 - verify that the collected information is consistent and correct
 - verify also with examples (related to general and borderline cases)
 - require precise definitions and classifications
 - distinguish the essential requirements from nice-to-have and less stringent requirements
 - proceed through subsequent refinements



8

8

Requirements analysis

- In practice:
 - choose the right abstraction level
 - standardize the structure of the sentences
 - avoid convoluted sentences
 - identify synonyms/homonyms and unify the terms
 - make explicit all references between terms
 - build a glossary of terms



9

9

Conceptual design

- Various project strategies have been proposed
- The most effective strategy is hybrid
 - the basic concepts are identified (the most relevant entities and relationships)
 - the initial project is progressively refined by adding the attributes, the cardinality of the relationships, the hierarchies, other entities and relationships
- If the problem is highly complex, it can be broken down into subproblems to solve separately and integrate later



10

10

Conceptual design: general criteria

- If a concept has significant properties or describes classes of objects with autonomous existence
 - entity
- If a concept has a simple structure and has no relevant properties
 - attribute (possibly multivalued)
- If two or more concepts are related
 - relation
- If a concept is a particular case of another one
 - hierarchy



11

11

Quality of a conceptual scheme

- **Correctness**
 - use of appropriate model constructs
 - absence of syntactic and semantic errors
- **Completeness**
 - representation of all the concepts of interest
- **Minimality**
 - every requirement is represented only once in the schema
 - all redundancies are verified and documented
- **Legibility**



12

12

Database design

Problem requirements



13

13

Problem requirements

- We want to represent a database for the management of a medical examination booking system within a Local Health Authority (ASL), considering the following information.
- Each patient is characterized by the health card number, name, surname, address, date of birth, place of birth and age.
- ASL hospitals are characterized by a numerical code, name and address.



14

14

Problem requirements

- Each hospital is divided into departments identified by a unique numeric code within the hospital and characterized by department name and telephone number.
- The department staff is identified by a Social Security Number (SSN). Name, surname and address are also known.
- Among the staff, for doctors the list of specializations achieved is known and for the volunteers the name of the association they belong to (if available) is stored.



15

15

Problem requirements

- The medical examinations that can be performed are characterized by a numerical code and a textual description (e.g., X-ray exam, etc.).
- For specialistic examinations, the doctor who carries out the visit and a description of the diet to follow (if necessary) are also stored.
- The laboratories that perform the examinations are identified by a unique code within the hospital and they are characterized by the name of the laboratory, the location plan and the room number.



16

16

Problem requirements

- For each member of the laboratory staff, the days and laboratories in which he/she works are stored. Pay attention to the fact that during the same day each staff member can work in more than one laboratory.



17

17

Problem requirements

- Each exam requires a reservation. For each exam reservation made by a patient, the date and time of the exam, the laboratory where it is performed, the cost of the ticket and the information about the exam being urgent or not are stored.
- Please note that each patient can make multiple reservations for the same exam on different dates and the same exam cannot be repeated on the same day by the same patient, even in different laboratories.

DBG

18

18

Problem requirements

- Each doctor can take on different roles during his/her career (e.g. assistant, head physician, etc.). We want to keep track of the roles each doctor has taken on during his/her career and the related time periods (start date, end date).
- Keep in mind that each doctor cannot take on more than one role at the same time, but he/she can take on the same role in different time periods.

DBG

19

19

Database design

Main concepts

DBG

20

20

Main concepts identification

- Text analysis to identify the most important concepts
 - the main entities of the ER diagram
 - any links between entities

DBG

21

21

Patient Concept

- Each *patient* is characterized by health card number, name, surname, address, date of birth, place of birth and age.

DBG

22

22

Patient Concept

Patient

DBG

23

23

Hospital concept

- ASL *hospitals* are characterized by a numeric code, a name and an address.



24

24

Hospital concept

Patient

Hospital



25

25

Department concept

- Each hospital is divided into *departments* identified by a unique numeric code within the hospital and characterized by the department name and the telephone number.



26

26

Department concept

Patient

Hospital

Department



27

27

Staff concept

- The *department staff* is identified by a Social Security Number (SSN).
- Name, surname and address are also known.
- Among the staff, for *doctors* the list of specializations achieved is known and for the *volunteers* the name of the association they belong to (if available) is stored.



28

28

Staff concept

Patient

Staff

Hospital

Department



29

29

Examination concept


- The *medical examinations* that can be performed are characterized by a numerical code and a textual description (e.g. X-ray exam, etc.).
- For specialist examinations, the doctor who carries out the visit and a description of the diet to follow (if necessary) are also stored.

D3G

30

30

Examination concept



D3G

31

31

Laboratory concept


- The *laboratories* that perform the tests are identified by a unique code within the hospital and are characterized by the name of the laboratory, the location plan and the room number.

D3G

32

32

Laboratory concept




D3G

33

33

Main concepts

- Main concepts
 - patient
 - examination
 - laboratory
 - hospital
 - department
 - staff



D3G

34

34

Database design

Model refinement (I)

D3G

35

35

Concepts refinement

- Refinement of concepts
 - introduction of hierarchies
 - attributes definition
 - characterization of relationships with cardinality

36

36


Staff hierarchy

- The department staff is identified by a Social Security Number (SSN). Name, surname and home address are also known.
- Among the staff, for *doctors* the list of specializations achieved is known and for the *volunteers* the name of the association they belong to (if available) is stored.

37

37

Staff hierarchy



38

38


Exam hierarchy

- The medical examinations that can be performed are characterized by a numerical code and a textual description (e.g. radiography, etc.). For *specialist examinations*, the doctor who carries out the visit and a description of the diet to follow (if necessary) are also stored.

39

39


Exam hierarchy




40

40

Refinement of the Patient entity



Each patient is characterized by health card number, name, surname, address, date of birth, place of birth and age.



42

42

Date of birth and Age Attributes


- The Age attribute is redundant because it can be easily calculated starting from the date of birth (BirthDate)
- This information must be attached to the conceptual model documentation
 - Age derivation rule from BirthDate: Age = Year (Today ()) - BirthDate)
- Elimination of the Age attribute will be evaluated during the restructuring phase of the ER scheme

DBGi

43

43

Refinement of the Hospital entity



DBGi

45

45

Database design

Model refinement (II)

DBGi

46

46

Relationship between staff and department


- The **department staff** is identified by a Social Security Number (SSN). Name, surname and address are also known. Among the staff, for doctors the list of specializations achieved is known and for volunteers the name of the association they belong to (if available) is stored.

DBGi

47

47

Relationship between staff and department




DBGi

48

48

Cardinality of the work in Relationship




DBGi

49

49

Cardinality of the *work in* Relationship




D&G

50

50

Refinement of the Staff entity




D&G

51

51

Refinement of the Doctor and Volunteer entities




D&G

53

53

Refinement of the Doctor and Volunteer entities




D&G

54

54

Refinement of the Examination entity




D&G

56

56

Refinement of the Specialistic Examination entity




D&G

58


58

Refinement of the Specialistic Examination entity




59

Cardinality of the *Carried out by* relationship



60

Cardinality of the *Carried out by* relationship



61

Database design

Model refinement (III)


62

Relationship between Department and Hospital

- Each hospital is divided into departments identified by a unique numeric code within the hospital and characterized by the department name and the telephone number.


63

Relationship between Department and Hospital




64

Cardinality of the *Hosts* relationship



65

Cardinality of the *Hosts* relationship




66

Department entity identifier

- Each hospital is divided into departments identified by a unique numeric code within the hospital and characterized by the department name and the telephone number.


67

Department entity identifier



68

Refinement of the Department entity



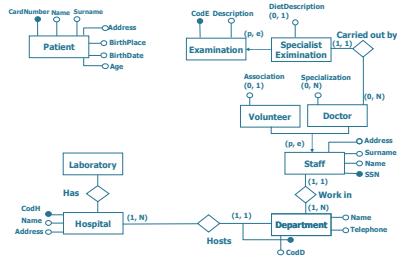
69

Relationship between Laboratory and Hospital

- The laboratories that perform the tests are identified by a unique code within the hospital and are characterized by the name of the laboratory, the location plan and the room number.

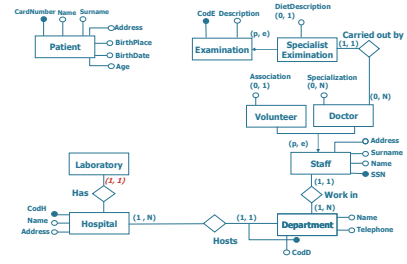
70

Relationship between Laboratory and Hospital



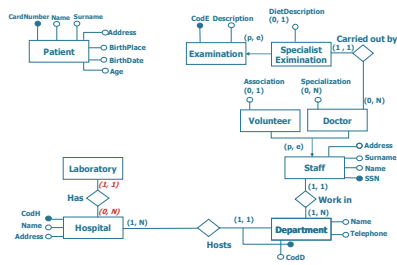
71

Cardinality of the Has relationship



72

Cardinality of the Has relationship



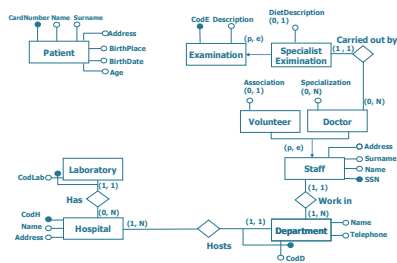
73

Relationship between Laboratory and Hospital

- The laboratories that perform the tests are identified by a unique code within the hospital and they are characterized by the name of the laboratory, the location plan and the room number.

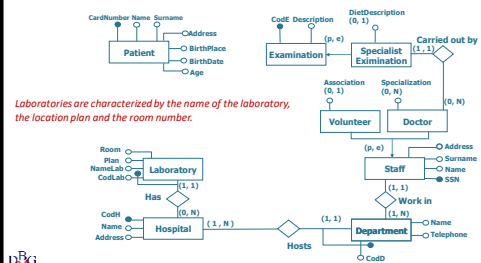
74

Laboratory entity identifier



75

Refinement of the Laboratory entity



76

Database design

Representation of time (I)

DBG

77

77

Time representation

- It is necessary to explicitly represent time-related information in the case of
 - representation of events
 - changes in the information content of entities or time attributes
- Various patterns are possible
 - by means of N-ary relationship with a time entity
 - through historicized entities
 - through binary relationships with a time entity

DBG

78

78

Relationship between Staff and Laboratory

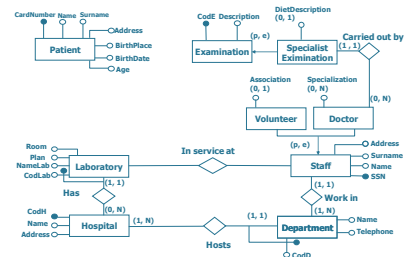
- For each member of the laboratory staff, the days and laboratories in which he/she works are stored. It should be borne in mind that during the same day each member of the staff can work in several laboratories.

DBG

79

79

Relationship between Staff and Laboratory




DBG

80

80

Relationship between Staff and Laboratory




DBG

81

81

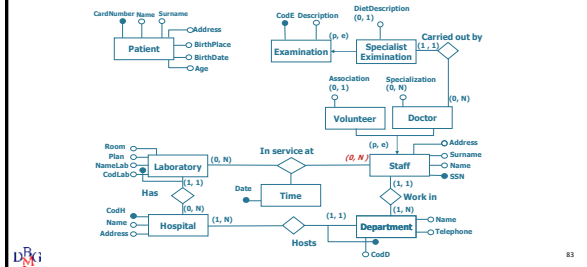
Cardinality of the *In service at* relationship



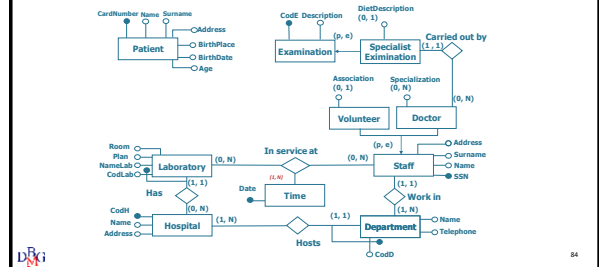
DBG

82

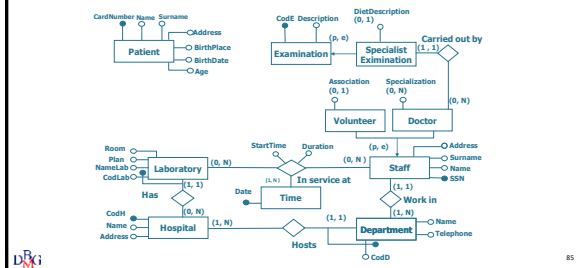
82

Cardinality of the *In service at* relationship

83

Cardinality of the *In service at* relationship

84

Cardinality of the *In service at* relationship

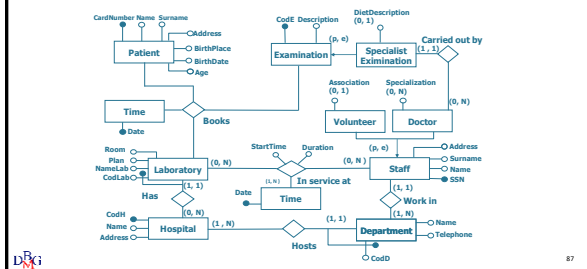
85

Representation of the reservation

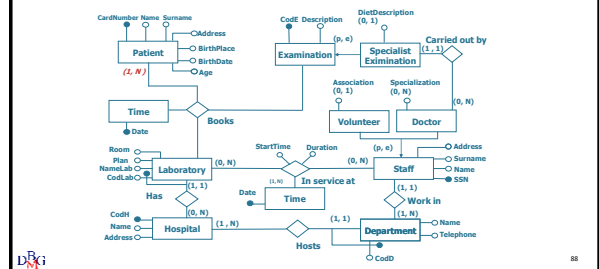
- Each exam requires a reservation.
- For each exam reservation made by a patient, the date and time of the exam, the laboratory where it is performed, the cost of the ticket and the information about the exam being urgent or not are stored.
- Please note that each patient can make multiple reservations for the same exam on different dates and the same exam cannot be repeated on the same day by the same patient, even in different laboratories.

86

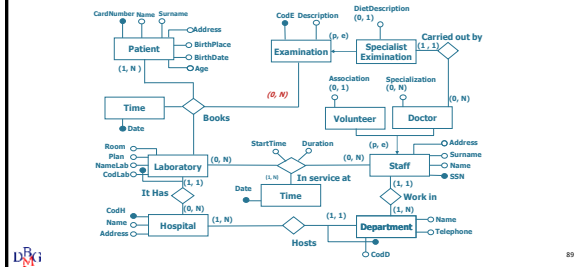
Representation of the reservation



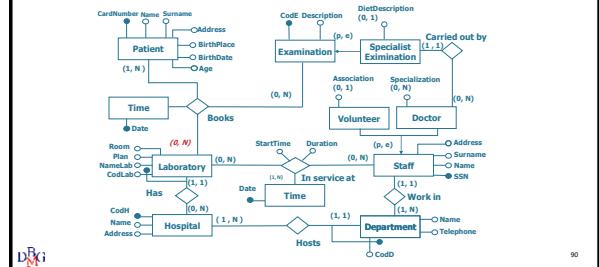
87

Cardinality of the *Books* relationship

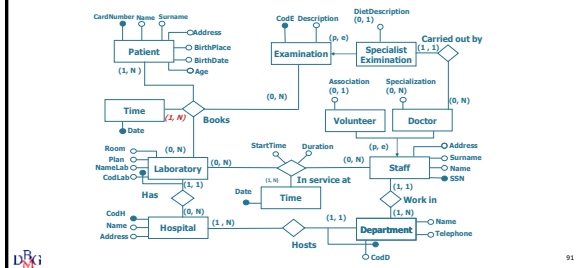
88

Cardinality of the *Books* relationship

89

Cardinality of the *Books* relationship

90

Cardinality of the *Books* relationship

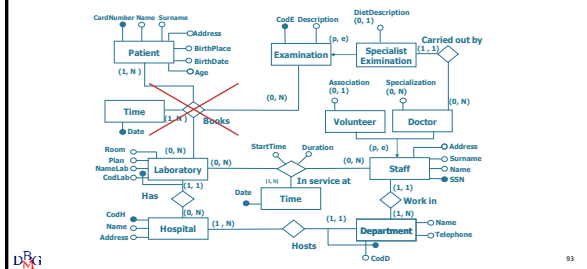
91

Reservation constraints

- Each exam requires a reservation.
- For each exam reservation made by a patient, the date and time of the exam, the laboratory where it is performed, the cost of the ticket and the information about the exam being urgent or not are stored.
- Please note that each patient can make multiple reservations for the same exam on different dates and the same exam cannot be repeated on the same day by the same patient, even in different laboratories.*

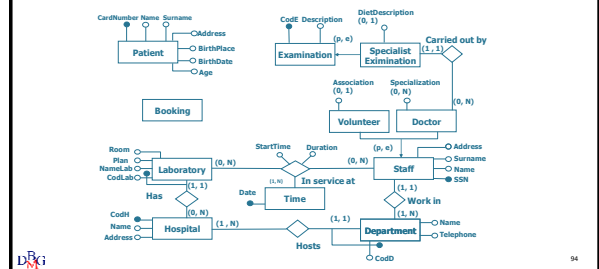
92

Representation of the reservation



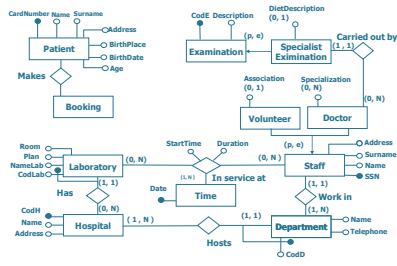
93

Introducing the Booking entity



94

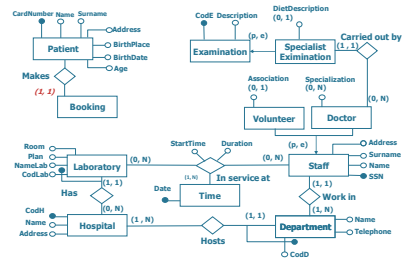
Relationship between Booking and Patient



D&G

95

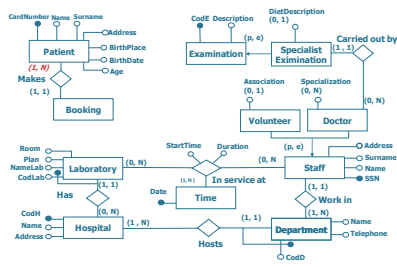
95

Cardinality of the *Makes* relationship

D&G

96

96

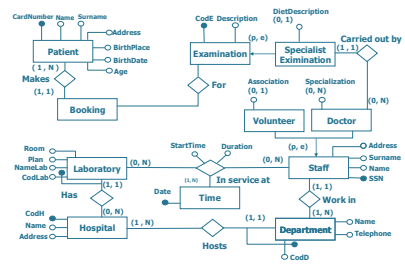
Cardinality of the *Makes* relationship

D&G

97

97

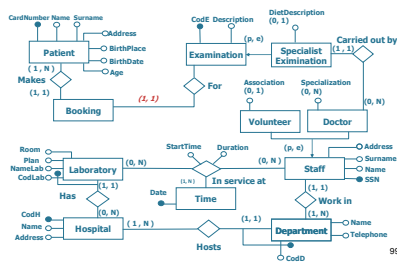
Relationship between Booking and Patient



D&G

98

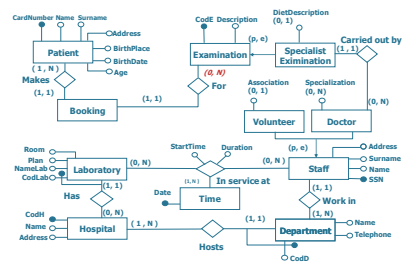
98

Cardinality of the *For* relationship

D&G

99

99

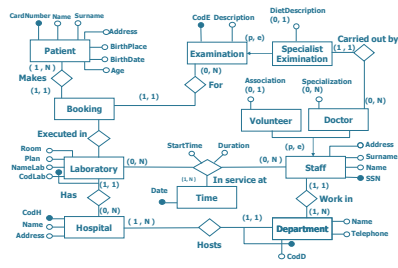
Cardinality of the *For* relationship

D&G

100

100

Relationship between Booking and Laboratory

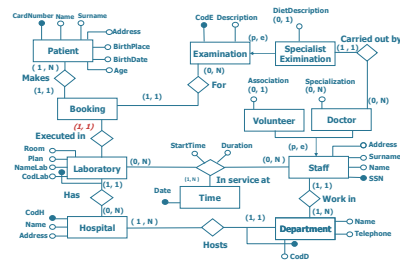


D&G

101

101

Cardinality of the *Executed in* relationship

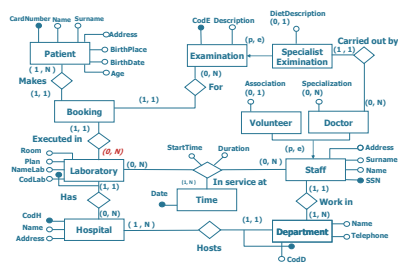


D&G

102

102

Cardinality of the *Executed in* relationship

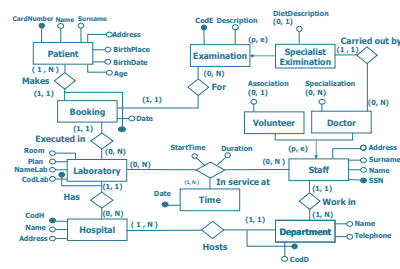


D&G

103

103

Booking entity identifier

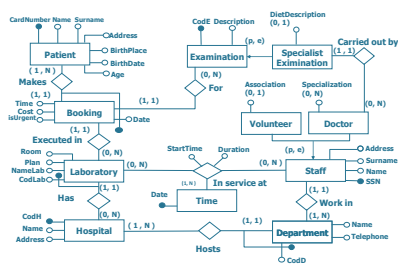


D&G

104

104

Refinement of the Booking entity



D&G

105

105

Relationship between Doctor and Role

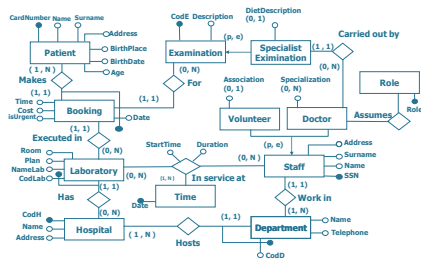
- Each doctor can take on different roles during his/her career (e.g. assistant, head physician, etc.). We want to keep track of the roles each doctor has taken on during his/her career and the related time periods (start date, end date). Keep in mind that each doctor cannot take on more than one role at the same time, but he/she can take on the same role in different time periods.

D&G

106

106

Relationship between Doctor and Role

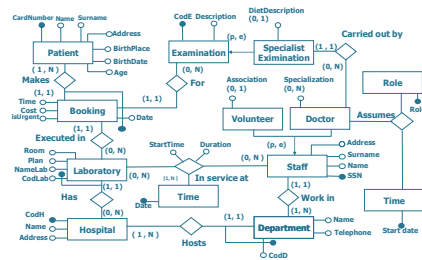


D&G

107

107

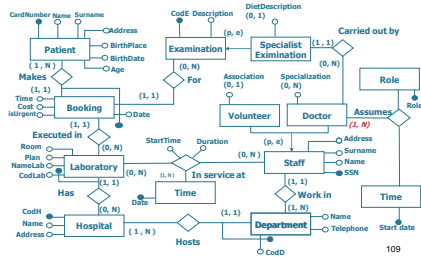
Relationship between Doctor and Role



D&G

108

108

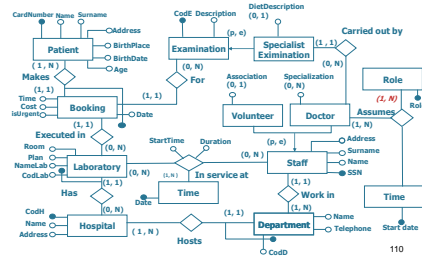


D&G

109

109

Cardinality of the Assumes relationship

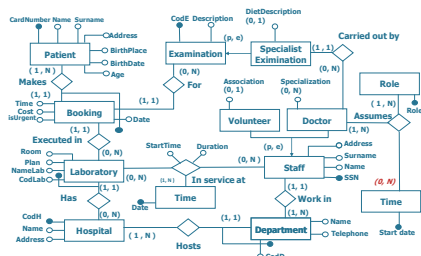


D&G

110

110

Cardinality of the Assumes relationship

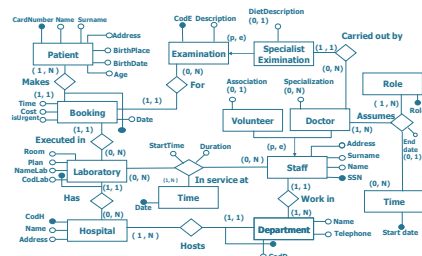


D&G

111

111

Cardinality of the Assumes relationship



D&G

112

112

Constraints on the *Assumes* relationship


- Each doctor can take on different roles during his/her career (e.g. assistant, head physician, etc.). We want to keep track of the roles each doctor has taken on during his/her career and the related time periods (start date, end date).
Keep in mind that each doctor cannot take on more than one role at the same time, but he/she can take on the same role in different time periods.

DBG

113

113

Historicization of the role of the doctor




DBG

114

114

Historicization of the role of the doctor




DBG

115

115

Historicization of the role of the doctor




DBG

116

116

Historicization of the role of the doctor




DBG

117

117

Historicization of the role of the doctor

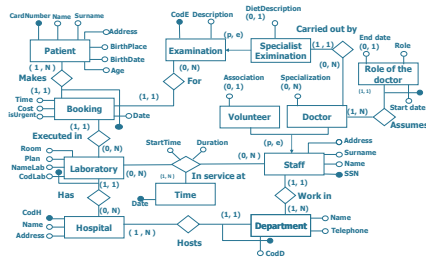


DBG

118

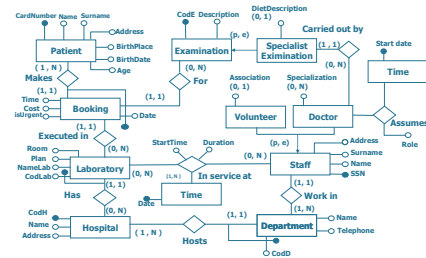
118

Historicization of the role of the doctor



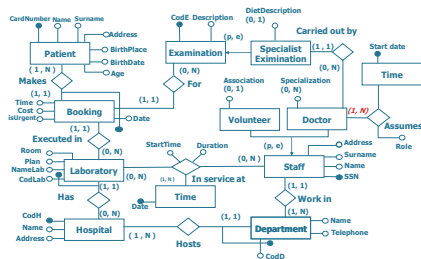
119

Role history: alternative modeling



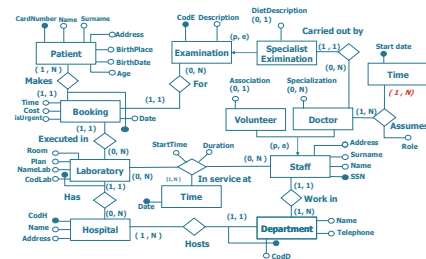
120

Cardinality of the Assumes relationship



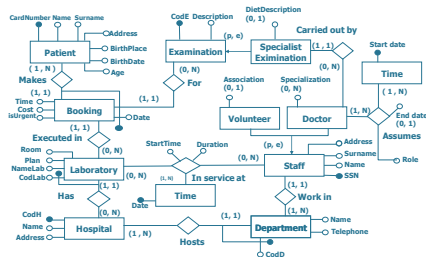
121

Cardinality of the Assumes relationship



122

Cardinality of the Assumes relationship



123

Politecnico di Torino

DBG

Logic design

Example

0

Example of relational logic design

- Introduction
- ER schema restructuring
- Translation of the entities without an external identifier
- Translation of the entities with an external identifier
- Translation of the relationships

DBG

1

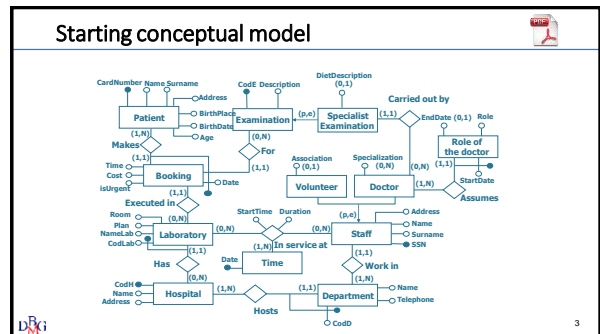
Politecnico di Torino

DBG

Introduction

Example of relational logic design

2



3

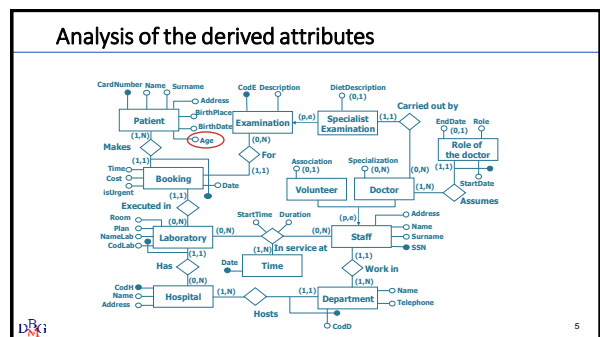
Politecnico di Torino

DBG

Restructuring the ER model


Example of relational logic design

4



5

Derived Age attribute




- The Age attribute can be removed since
 - it can be easily calculated from the date of birth (BirthDate)
 - it is not generally present in a query

D&G

6

6

Eliminating the Age attribute

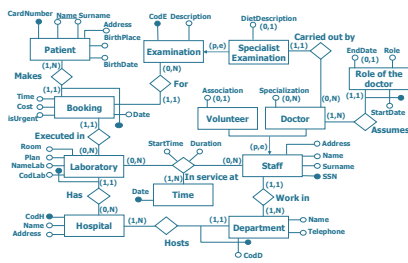


D&G

7

7

Restructured schema (n. 1)

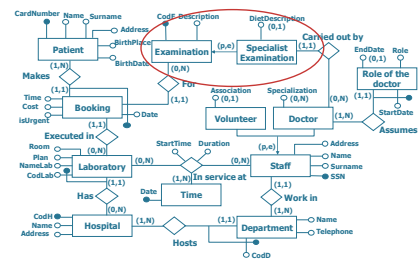


D&G

8

8

Removing generalizations




D&G

9

9

Removing generalizations




D&G

10

10

Merge into the parent entity




D&G

11

11

Restructured schema (n.2)




D&G

12

12

Removing generalizations




D&G

13

13

Staff generalization




D&G

14

14

Merge into the parent entity




D&G

15

15

Restructured schema (n.3)




D&G

16

16

Removing multivalued attributes



D&G

17

17

Removing multivalued Specialization attribute



DBG

18

18

New specialization entity



DBG

19

19

Cardinality of the Has relationship

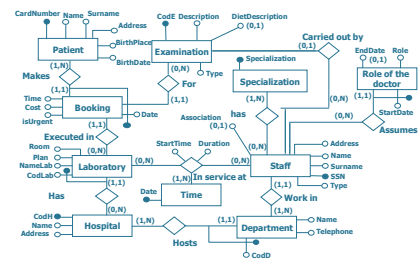


DBG

20

20

Restructured schema (final)



DBG

21

21



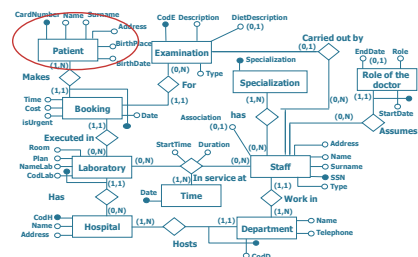
DBG

Traslating entities without an external identifier

Example of relational logic design

22

Translation of the Patient entity



DBG

23

23

Translation of the Patient entity


Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)



24

24

Translation of the Hospital entity



25

25

Translation of the Hospital entity

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)

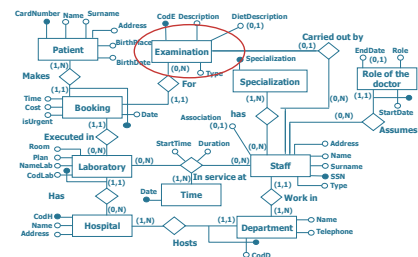
Hospital(CodH, Name, Address)



26

26

Translation of the Examination entity



27

27

Translation of the Examination entity

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)

Hospital(CodH, Name, Address)


Examination(CodE, Description, DietDescription*, Type)



28

28

Translation of the Staff entity



29

29

Translation of the Staff entity

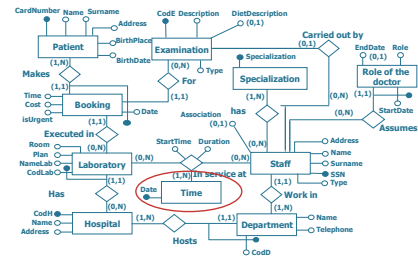
Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)



30

30

Translation of the Time entity




31

31

Translation of the Time entity

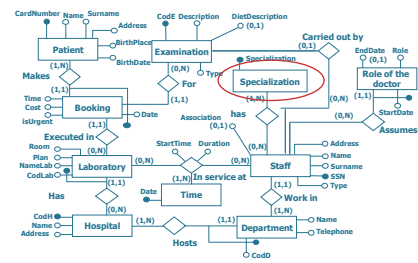
Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)



32

32

Translation of the Specialization entity




33

33



Translation of the Specialization entity

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)
 Specialization(Specialization)



34

34





Translating entities with an external identifier

Example of relational logic design

35

Translation of the Laboratory entity



D8gi

36

36

Translation of the Laboratory entity

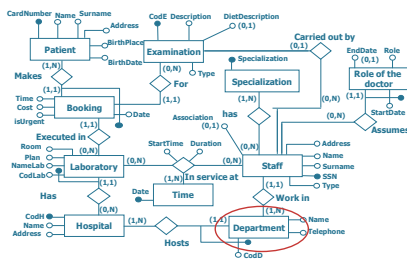
Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(Code, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)

D8gi

37

37

Translation of the Department entity



D8gi

38

38

Translation of the Department entity


Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(Code, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)

D8gi

39

39

Translation of the Booking entity



D8gi

40

40

Translation of the Booking entity

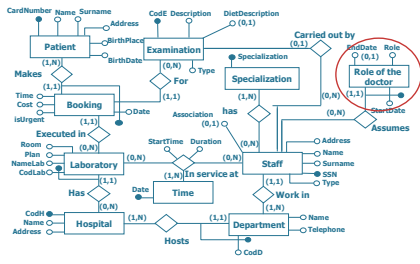
Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(Code, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, Code, Date, Time, Cost, isUrgent)

D8gi

41

41

Translation of the Role of the doctor entity



DBG

42

42

Translation of the Role of the doctor entity

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent)
 DoctorRole(SSN, StartDate, EndDate*, Role)

DBG

43

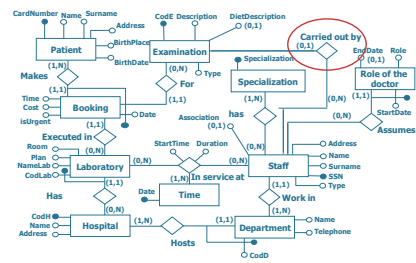
43

Translating relationships

Example of relational logic design

44

Binary one-to-many *Carried out by* relationship



DBG

45

45

Binary one-to-many *Carried out by* relationship

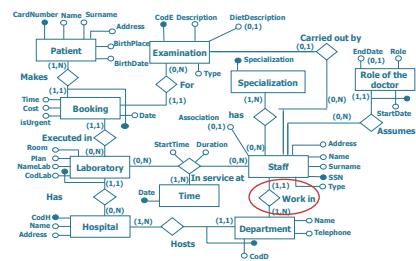
Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent)
 DoctorRole(SSN, StartDate, EndDate*, Role)

DBG

46

46

Binary one-to-many *Work in* relationship



DBG

47

47

Binary one-to-many *Work in* relationship


Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type, *CodD*, *CodH*)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent)
 DoctorRole(SSN, StartDate, EndDate*, Role)

D8gi

48

48

Binary one-to-many *Executed in* relationship



D8gi

49

49

Binary one-to-many *Executed in* relationship


Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, *CodLab*, *CodH*)
 DoctorRole(SSN, StartDate, EndDate*, Role)

D8gi

50

50

Binary many-to-many *has* relationship



D8gi

51

51

Binary many-to-many *has* relationship


Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, CodLab, CodH)
 DoctorRole(SSN, StartDate, EndDate*, Role)
HasSpecialization(SSN, Specialization)

D8gi

52

52

Ternary many-to-many *In service at* relationship



D8gi

53

53

Ternary many-to-many *In service at* relationship

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)
 Time(Date)
 Specialization(Specialization)
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, CodLab, CodH)
 DoctorRole(SSN, StartDate, EndDate*, Role)
 HasSpecialization(SSN, Specialization)
 InServiceAt(SSN, CodD, Date, StartTime, Duration)



Eliminating redundant tables

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)
 Time(Date)
~~Specialization(Specialization)~~
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, CodLab, CodH)
 DoctorRole(SSN, StartDate, EndDate*, Role)
 HasSpecialization(SSN, Specialization)
 InServiceAt(SSN, CodLab, CodD, Date, StartTime, Duration)



Eliminating redundant tables

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
 Hospital(CodH, Name, Address)
 Examination(CodE, Description, DietDescription*, Type, SSN*)
 Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)
~~Time(Date)~~
 Laboratory(CodLab, CodH, NameLab, Plan, Room)
 Department(CodD, CodH, Name, Telephone)
 Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, CodLab, CodH)
 DoctorRole(SSN, StartDate, EndDate*, Role)
 HasSpecialization(SSN, Specialization)
 InServiceAt(SSN, CodLab, CodD, Date, StartTime, Duration)




Final relational schema

Patient(CardNumber, Name, Surname, Address, BirthPlace, BirthDate)
Hospital(CodH, Name, Address)
Examination(CodE, Description, DietDescription*, Type, SSN*)
Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)
Laboratory(CodLab, CodH, NameLab, Plan, Room)
Department(CodD, CodH, Name, Telephone)
Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, CodLab, CodH)
DoctorRole(SSN, StartDate, EndDate*, Role)
HasSpecialization(SSN, Specialization)
InServiceAt(SSN, CodLab, CodD, Date, StartTime, Duration)



Referential integrity constraints

Example of relational logic design



Referential integrity: *Carried out by* relationship

- Involved tables

Examination(CodE, Description, DietDescription*, Type, **SSN***)

Staff(**SSN**, Name, Surname, Address, Association*, Type, CodD, CodH)

- Referential integrity constraint

Examination(SSN) REFERENCES Staff(SSN)



60

60

Referential integrity: *Work in* relationship

- Involved tables

Staff(**SSN**, Name, Surname, Address, Association*, Type, **CodD**, **CodH**)

Department(**CodD**, **CodH**, Name, Telephone)

- Referential integrity constraint

Staff(CodD, CodH) REFERENCES Department(CodD, CodH)



61

61

Referential integrity: *Hosts* relationship

- Involved tables

Department(**CodD**, **CodH**, Name, Telephone)

Hospital(**CodH**, Name, Address)

- Referential integrity constraint

Department(CodH) REFERENCES Hospital(CodH)



62

62

Referential integrity: *Has* relationship

- Involved tables

Laboratory(**CodLab**, **CodH**, NameLab, Plan, Room)

Hospital(**CodH**, Name, Address)

- Referential integrity constraint

Laboratory(CodH) REFERENCES Hospital(CodH)



63

63

Referential integrity: *Makes* relationship

- Involved tables

Booking(**CardNumber**, **Code**, **Date**, Time, Cost, isUrgent, CodLab, CodH)

Patient(**CardNumber**, Name, Surname, Address, BirthPlace, BirthDate)

- Referential integrity constraint

Booking(CardNumber) REFERENCES Patient(CardNumber)



64

64

Referential integrity: *For* relationship

- Involved tables

Booking(**CardNumber**, **Code**, **Date**, Time, Cost, isUrgent, CodLab, CodH)

Examination(**Code**, Description, DietDescription*, Type, **SSN***)

- Referential integrity constraint

Booking(Code) REFERENCES Examination(Code)



65

65

Referential integrity: *Executed in* relationship

- Involved tables

Booking(CardNumber, CodE, Date, Time, Cost, isUrgent, CodLab, CodH)
Laboratory(CodLab, CodH, NameLab, Plan, Room)

- Referential integrity constraint

Booking(CodLab, CodH) REFERENCES Laboratory(CodLab, CodH)



66

66

Referential integrity: *Assumes* relationship

- Involved tables

DoctorRole(SSN, StartDate, EndDate*, Role)
Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)

- Referential integrity constraint

DoctorRole(SSN) REFERENCES Staff(SSN)



67

67

Referential integrity: *Has* relationship

- Involved tables

HasSpecialization(SSN, Specialization)
Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)

- Referential integrity constraint

HasSpecialization(SSN) REFERENCES Staff(SSN)



68

68

Referential integrity: *In service at* relationship

- Involved tables

InServiceAt(SSN, CodLab, CodH, Date, StartTime, Duration)
Staff(SSN, Name, Surname, Address, Association*, Type, CodD, CodH)

- Referential integrity constraint

InServiceAt(SSN) REFERENCES Staff(SSN)



69

69

Referential integrity: *In service at* relationship

- Involved tables

InServiceAt(SSN, CodLab, CodH, Date, StartTime, Duration)
Laboratory(CodLab, CodH, NameLab, Plan, Room)

- Referential integrity constraint

InServiceAt(CodLab, CodH) REFERENCES Laboratory(CodLab, CodH)



70

70


Referential integrity constraints

Examination(SSN) REFERENCES Staff(SSN)
Staff(CodD, CodH) REFERENCES Department(CodD, CodH)
Department(CodH) REFERENCES Hospital(CodH)
Laboratory(CodH) REFERENCES Hospital(CodH)
Booking(CardNumber) REFERENCES Patient(CardNumber)
Booking(CodE) REFERENCES Examination(CodE)
Booking(CodLab, CodH) REFERENCES Laboratory(CodLab, CodH)
DoctorRole(SSN) REFERENCES Staff(SSN)
HasSpecialization(SSN) REFERENCES Staff(SSN)
InServiceAt(SSN) REFERENCES Staff(SSN)
InServiceAt(CodLab, CodH) REFERENCES Laboratory(CodLab, CodH)




71

71



Politecnico di Torino




Normalization

Database design


0

Normalization

- Introduction
- Normal form of Boyce Codd
- Decomposition in normal form
- Properties of decompositions
- Lossless decomposition
- Conservation of dependencies



1




Introduction

Normalization

2

Normalization

- Normalization is a process which, starting from a non-normalized relational schema, obtains a normalized relational schema
- Normalization is **not a design methodology**, but a **verification tool**
- The design methodology based on ER schemas normally produces normalized relational schemas
- Normalization checks can also be applied to ER schemas




3

Example

Exam Passed

StudentID	Residence	CourseID	CourseName	Grade
s94539	Milan	04FLYCY	Electronic calculators	30
s94540	Turin	01FLTCY	Database design	26
s94540	Turin	01KPNCY	Computer network	28
s94541	Pescara	01KPNCY	Computer network	29
s94542	Lecce	04FLYCY	Electronic calculators	25

- **Constraints**
 - The primary key is the pair StudentID, CourseID
 - The place of residence of each student is unique and is an attribute of the student alone, regardless of the exams he or she has passed
 - The name of the course is unique and is a function of the course only, regardless of which students pass the corresponding exam




4

Example: Redundancy

Exam Passed

StudentID	Residence	CourseID	CourseName	Grade
s94539	Milan	04FLYCY	Electronic calculators	30
s94540	Turin	01FLTCY	Database design	26
s94540	Turin	01KPNCY	Computer network	28
s94541	Pescara	01KPNCY	Computer network	29
s94542	Lecce	04FLYCY	Electronic calculators	25

- **Redundancy**
 - In all rows where a student appears, his or her place of residence is repeated
 - In all rows where the same course appears, its name is repeated



5

Example: Anomalies

Exam Passed

StudentID	Residence	CourseID	CourseName	Grade
s94539	Milan	04FLYCY	Electronic calculators	30
s94540	Turin	01FLTCY	Database design	26
s94540	Turin	01KPNCY	Computer network	28
s94541	Pescara	01KPNCY	Computer network	29
s94542	Lecce	04FLYCY	Electronic calculators	25

- **Update anomaly**
 - If a student's place of residence changes, all the rows in which it appears must be modified at the same time
- **Insertion anomaly**
 - If a new student enrolls at university, he or she cannot be entered in the database until he or she passes the first exam
- **Deletion anomaly**
 - If a student withdraws from studies, it is not possible to keep track of his place of residence

DBGi

6

6

Redundancy

- A single relation is used to represent heterogeneous information
 - some data are repeated in different tuples without adding new information
 - redundant data

DBGi

7

7

Anomalies

- Redundant information must be updated atomically (all at the same time)
- The deletion of a tuple implies the deletion of all concepts represented in it
 - including those that might still be valid
- The insertion of a new tuple is only possible if at least the complete information about the primary key exists
 - it is not possible to insert the part of the tuple relating to only one concept

DBGi

8

8

Boyce-Codd normal form

Normalization

DBGi

9

9

Functional dependency

- It is a special type of integrity constraint
- It describes functional links between the attributes of a relation
- Example: the place of residence is unique for each student
 - each time the same student appears, the value is repeated
 - the value of StudentID **determines** the value of Residence

Exam Passed

StudentID	Residence	CourseID	CourseName	Grade
s94539	Milan	04FLYCY	Electronic calculators	30
s94540	Turin	01FLTCY	Database design	26
s94540	Turin	01KPNCY	Computer network	28
s94541	Pescara	01KPNCY	Computer network	29
s94542	Lecce	04FLYCY	Electronic calculators	25

DBGi

10

10

Functional dependency

- A relation r satisfies the functional dependency $X \rightarrow Y$ if, for each pair t_1, t_2 of tuples of r , having the same values for attributes in X , t_1 and t_2 also have the same values for attributes in Y
 - X determines Y (in r)
- Examples

StudentID \rightarrow Residence
 CourseID \rightarrow CourseName
 StudentID CourseID \rightarrow CourseName

DBGi

11

11

Non-trivial dependency

- The dependency $\text{StudentID CourseID} \rightarrow \text{CourseID}$ is trivial because CourseID is part of both sides
- A functional dependency $X \rightarrow Y$ is non-trivial if no attribute in X appears among the attributes in Y



12

12

Functional dependencies and keys

- Given a key K of a relation r
 $K \rightarrow$ any other attribute of r (or set of attributes)
- Examples
 - $\text{StudentID CourseID} \rightarrow \text{Residence}$
 - $\text{StudentID CourseID} \rightarrow \text{CourseName}$
 - $\text{StudentID CourseID} \rightarrow \text{Grade}$



13

13

Functional dependencies and anomalies

- Anomalies** are caused by attribute properties involved in functional dependencies
 - Examples
 - $\text{StudentID} \rightarrow \text{Residence}$
 - $\text{CourseID} \rightarrow \text{CourseName}$
- Functional dependencies** on keys do not give rise to anomalies
 - Example
 - $\text{StudentID CourseID} \rightarrow \text{Grade}$



14

14

Functional dependencies and anomalies

- The anomalies are caused by
 - the inclusion of mutually independent concepts in the same relation
 - functional dependencies $X \rightarrow Y$ allowing for multiple tuples with the same value of X
 - X does not contain a key



15

15

Boyce Codd normal form (BCNF)

- BCNF = Boyce Codd Normal Form
- A relation r is in BCNF if, for every (non-trivial) functional dependency $X \rightarrow Y$ defined on it, X contains a key of r (X is superkey of r)
- Anomalies and redundancies are not present in BCNF relations because independent concepts are separated in different relations



16

16

Normal form decomposition

Normalization



17

17

BCNF decomposition

- Normalization
 - process of replacing a non-normalised relation by two or more relations in BCNF
- Criteria
 - a relation representing several independent concepts is decomposed into smaller relations, one for each concept, by means of functional dependencies
- The new relations are obtained by projections onto the sets of attributes corresponding to the functional dependencies
- The keys of the new relations are the left parts of the functional dependencies
 - the new relations are in BCNF

DBGi

18

18

Example

- Functional dependencies in the example
 - $\text{StudentID} \rightarrow \text{Residence}$
 - $\text{CourseID} \rightarrow \text{CourseName}$
 - $\text{StudentID CourseID} \rightarrow \text{Grade}$

Exam Passed

StudentID	Residence	CourseID	CourseName	Grade
s94539	Milan	04FLYCY	Electronic calculators	30
s94540	Turin	01FLTCY	Database design	26
s94540	Turin	01KPNCY	Computer network	28
s94541	Pescara	01KPNCY	Computer network	29
s94542	Lecce	04FLYCY	Electronic calculators	25

DBGi

19

19

Example

- By
 $R(\text{StudentID}, \text{Residence}, \text{CourseID}, \text{CourseName}, \text{Grade})$
- Functional dependencies in the example
 - $\text{StudentID} \rightarrow \text{Residence}$
 - $\text{CourseID} \rightarrow \text{CourseName}$
 - $\text{StudentID CourseID} \rightarrow \text{Grade}$
- The relations in BCNF are
 - $R_1(\text{StudentID}, \text{Residence}) = \pi_{\text{StudentID}, \text{Residence}} R$
 - $R_2(\text{CourseID}, \text{CourseName}) = \pi_{\text{CourseID}, \text{CourseName}} R$
 - $R_3(\text{StudentID}, \text{CourseID}, \text{Grade}) = \pi_{\text{StudentID}, \text{CourseID}, \text{Grade}} R$

DBGi

20

20

Example

R_1		R_2	
StudentID	Residence	CourseID	CourseName
s94539	Milan	04FLYCY	Electronic calculators
s94540	Turin	01FLTCY	Database design
s94540	Turin	01KPNCY	Computer network
s94541	Pescara		


R_3		
StudentID	CourseID	Grade
s94539	04FLYCY	30
s94540	01FLTCY	26
s94540	01KPNCY	28
s94541	01KPNCY	29
s94542	04FLYCY	25

DBGi

21

21

Example: corresponding ER scheme



Student (StudentID, Residence)
 Course (CourseID, CourseName)
 Exam Passed (StudentID, CourseID, Grade)

DBGi

22

22

Decomposition properties

Normalization

DBGi

23

23

Decomposition properties

- Are all decompositions acceptable?
 - essential properties for "good" decomposition
- Problems
 - information loss
 - loss of dependencies

DBGI

24

24

Example

Employee	Category	Salary
Rossi	2	1800
Verdi	3	1800
Bianchi	4	2500
Neri	5	2500
Bruni	6	3500

R (Employee, Category, Salary)

Employee → Category
Employee → Salary
Category → Salary

DBGI

25

25

Lossless Decomposition

Normalization

DBGI

26

26

Example: decomposition (n.1)

R (Employee, Category, Salary)

- Decomposition based on functional dependencies
 - Employee → Salary
 - Category → Salary

DBGI

27

27

Example: decomposition (n.1)

R (Employee, Category, Salary)

- Decomposing

$R_1(\text{Employee, Salary}) = \pi_{\text{Employee, Salary}} R$

Employee	Salary
Rossi	1800
Verdi	1800
Bianchi	2500
Neri	2500
Bruni	3500

$R_2(\text{Category, Salary}) = \pi_{\text{Category, Salary}} R$

Category	Salary
2	1800
3	1800
4	2500
5	2500
6	3500

DBGI

28

28

Example: recomposition (n.1)

- Recomposing

$R_1 \bowtie R_2$

Employee	Category	Salary
Rossi	2	1800
Rossi	3	1800
Verdi	2	1800
Verdi	3	1800
Bianchi	4	2500
...

← "spurious" tuples

- Reconstruction with loss of information

DBGI

29

29

Decomposition without loss

- The decomposition of a relation r into two sets of attributes X_1 and X_2 is **lossless** if the join of the projections of r into X_1 and X_2 is equal to r itself (no "spurious" tuples)
- A decomposition performed to normalize a relation must be lossless

DBGI

30

30

Lossless decomposition

- Given the relation $r(X)$ and sets of attributes X_1 and X_2 such that

$$X = X_1 \cup X_2$$

$$X_0 = X_1 \cap X_2$$
 if r satisfies the functional dependency

$$X_0 \rightarrow X_1 \text{ or } X_0 \rightarrow X_2$$
 the decomposition of r on X_1 and X_2 is lossless
- Common attributes form a key to at least one of the decomposed relations

DBGI

31

31

Example: loss of information

 $R_1 (\text{Employee}, \text{Salary}) \quad R_2 (\text{Category}, \text{Salary})$

- Verification of condition for lossless decomposition

$$X_1 = \text{Employee, Salary}$$

$$X_2 = \text{Category, Salary}$$

$$X_0 = \text{Salary}$$

- The attribute Salary does not satisfy the condition for lossless decomposition

DBGI

32

32

Example: decomposition (n.2)

 $R (\text{Employee}, \text{Category}, \text{Salary})$

- Decomposition based on functional dependencies

$$\text{Employee} \rightarrow \text{Category}$$

$$\text{Employee} \rightarrow \text{Salary}$$

- Decomposing

 $R_1 (\text{Employee}, \text{Category}) =$
 $\pi_{\text{Employee, Salary}} R$

Employee	Category
Rossi	2
Verdi	3
Bianchi	4
Neri	4
Bruni	5

 $R_2 (\text{Employee}, \text{Salary}) =$
 $\pi_{\text{Category, Salary}} R$

Employee	Salary
Rossi	1800
Verdi	1800
Bianchi	2500
Neri	2500
Bruni	3500

DBGI

33

33

Example: lossless decomposition?

 $R_1 (\text{Employee}, \text{Category}) \quad R_2 (\text{Employee}, \text{Salary})$
 $R_1 \bowtie R_2$

- Is the decomposition **lossless**?
- Verifying the condition for lossless decomposition

$$X_1 = \text{Employee, Category}$$

$$X_2 = \text{Employee, Salary}$$

$$X_0 = \text{Employee}$$

- The attribute Employee satisfies the condition for lossless decomposition

DBGI

34

34

Conservation of dependencies

Normalization

DBGI

35

35

Example: inserting a new tuple

R_1 (Employee, Category) R_2 (Employee, Salary)

- Inserting the tuple
 - Employee: Gialli – Category: 3 – Salary: 3500

Employee	Category
Rossi	2
Verdi	3
Bianchi	4
Neri	4
Bruni	5

Gialli	3
--------	---

Employee	Salary
Rossi	1800
Verdi	1800
Bianchi	2500
Neri	2500
Bruni	3500

Gialli	3500
--------	------

DBGI

36

36

Example: inserting a new tuple

- What happens if we insert the tuple (Gialli, 3500) in R_2 ?
 - in the original relation insertion is forbidden because it violates the dependency $\text{Category} \rightarrow \text{Salary}$
 - in the decomposition it is no longer possible to detect the violation, since the attributes Category and Salary are in separate relations
- The dependency between Category and Salary has been lost

DBGI

37

37

Conservation of dependencies

- A decomposition preserves dependencies if each of the functional dependencies of the original schema is present in one of the decomposed relations
- Dependencies should be retained to ensure that the same constraints are satisfied in the decomposed schema as in the original schema

DBGI

38

38

Example: decomposition (n.3)

R (Employee, Category, Salary)

- Decomposition based on functional dependencies
 - $\text{Employee} \rightarrow \text{Category}$
 - $\text{Category} \rightarrow \text{Salary}$

- Decomposing

R_1 (Employee, Category) = $\pi_{\text{Employee, Category}} R$

Employee	Category
Rossi	2
Verdi	3
Bianchi	4
Neri	4
Bruni	5

R_2 (Category, Salary) = $\pi_{\text{Category, Salary}} R$

Category	Salary
2	1800
3	1800
4	2500
5	2500
6	3500

DBGI

39

39

Example: Lossless decomposition

- Recomposing
 - $R_1 \bowtie R_2$
- Condition check for **lossless** decomposition
 - $X_1 = \text{Employee, Category}$
 - $X_2 = \text{Category, Salary}$
 - $X_0 = \text{Category}$
- The attribute Category satisfies the condition for lossless decomposition

DBGI

40

40

Example: Conservation of functional dependencies


- Recomposing
 - $R_1 \bowtie R_2$
- Conserved functional dependencies
 - $\text{Employee} \rightarrow \text{Category}$
 - $\text{Category} \rightarrow \text{Salary}$
- Functional dependency
 - $\text{Employee} \rightarrow \text{Salary}$
 can be reconstructed from
 - $\text{Employee} \rightarrow \text{Category}$
 - $\text{Category} \rightarrow \text{Salary}$

DBGI

41

41

Example: corresponding ER scheme



Employee (Employee, Category)
 Category (Category, Salary)

DBGI

42

42

Quality of a decomposition

- Decompositions must always satisfy the properties
 - lossless decomposition
 - ensures that the information in the original relation is accurately reconstructed (without spurious tuples) from the information in the decomposed relations
 - conservation of dependencies
 - ensures that the decomposed relations have the same capacity as the original relation to represent the integrity constraints

DBGI

43

43